

Lecture 4


Model Order Selection

EE-UY 4563/EL-GY 9123: INTRODUCTION TO MACHINE LEARNING
PROF. SUNDEEP RANGAN (WITH MODIFICATION BY YAO WANG)

Learning Objectives

- ❑ Compute the model order for a given model class
- ❑ Visually identify overfitting and underfitting of a model in a scatterplot
- ❑ Determine if there is under-modeling for a given true function and model class
- ❑ Compute the bias and variance for linear models (advanced)
- ❑ Perform cross-validation for selecting an optimal order
- ❑ Use cross-validation for feature selection

Outline

 Motivating Example: What polynomial degree should a model use?

- ☐ Bias and variance
- ☐ Bias and variance for linear models (Advanced)
- ☐ Cross-validation
- ☐ Feature selection

Polynomial Fitting

□ Last lecture: polynomial regression

□ Given data $(x_i, y_i), i = 1, \dots, N$

□ Learn a polynomial relationship:

$$y = \beta_0 + \beta_1 x + \dots + \beta_d x^d + \epsilon$$


- d = degree of polynomial. Called **model order**

- $\boldsymbol{\beta} = (\beta_0, \dots, \beta_d)$ = coefficient vector

□ Given d , can find $\boldsymbol{\beta}$ via least squares


□ How do we select d from data?



□ This problem is called **model order selection**.



Demo on Github

□ Demo on github: https://github.com/sdrangan/introml/blob/master/unit04_model_sel/demo1_polyfit.ipynb

□  GitHub, Inc. [US] | https://github.com/sdrangan/introml/blob/master/model_sel/polyfit.ipynb

Suggested Sites  Web Slice Gallery  Import to Mendeley

Demo: Polynomial Model Order Selection

In this demo, we will illustrate the process of cross-validation for model order selection. We der data for a polynomial fit. The lab will demonstrate how to:

- Characterize the model order for a simple polynomial model
- Measure training and test error for a given model order
- Select a suitable model order using cross-validation
- Plot the results for the model order selection process

We first load the packages as usual.

```
In [2]: import numpy as np
import matplotlib
import matplotlib.pyplot as plt
from sklearn import datasets, linear_model, preprocessing
%matplotlib inline
```

Polynomial Data


To illustrate the concepts, we consider a simple polynomial model:

$$y = \beta_0 + \beta_1 x + \dots + \beta_d x^d + \epsilon,$$

where d is the polynomial degree. We first generate synthetic data for this model.

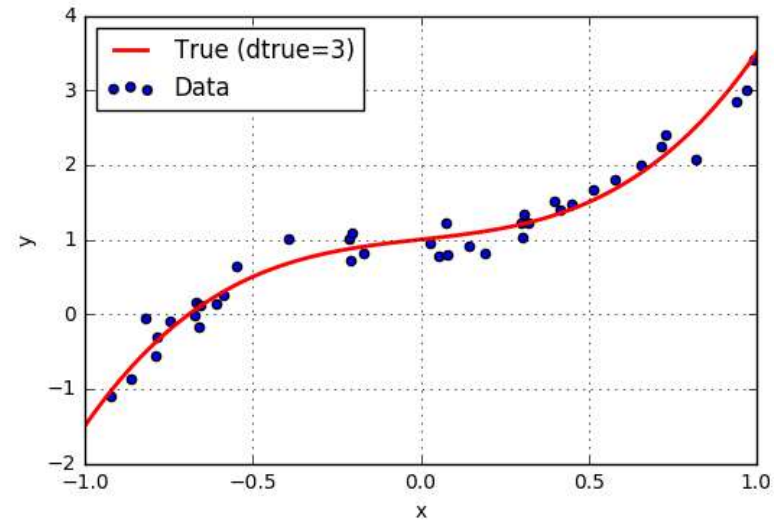
Example Question

- ❑ You are given some data.
- ❑ Want to fit a model: $y \approx f(x)$
- ❑ Decide to use a polynomial:
$$f(x) = \beta_0 + \beta_1 x + \cdots + \beta_d x^d$$
- ❑ What model order d should we use?
- ❑ Thoughts?



Synthetic Data

- Previous example is synthetic data
- x_i : 40 samples uniform in $[-1,1]$
- $y = f(x) + \epsilon$,
 - $f(x) = \beta_0 + \beta_1 x + \dots + \beta_d x^d = \text{"true relation"}$
 - $d = 3$, $\epsilon \sim N(0, \sigma^2)$
- Synthetic data useful for analysis
 - Know "ground truth"
 - Can measure performance of various estimators



```
# Import useful polynomial library
import numpy.polynomial.polynomial as poly


# True model parameters
beta = np.array([1,0.5,0,2]) # coefficients
wstd = 0.2 # noise
dtrue = len(beta)-1 # true poly degree

# Independent data
nsamp = 40
xdat = np.random.uniform(-1,1,nsamp)

# Polynomial
y0 = poly.polyval(xdat,beta)
ydat = y0 + np.random.normal(0,wstd,nsamp)
```

Fitting with True Model Order

- Suppose true polynomial order, $d=3$, is known
- Use linear regression
 - numpy.polynomial package
- Get very good fit




```
d = 3
beta_hat = poly.polyfit(xdat,ydat,d)

# Plot true and estimated function
xp = np.linspace(-1,1,100)
yp = poly.polyval(xp,beta)
yp_hat = poly.polyval(xp,beta_hat)
plt.xlim(-1,1)
plt.ylim(-1,3)
plt.plot(xp,yp,'r-',linewidth=2)
plt.plot(xp,yp_hat,'g-',linewidth=2)


# Plot data
plt.scatter(xdat,ydat)
plt.legend(['True (dtrue=3)', 'Est (d=3)', 'Data'], loc='upper left')
plt.grid()
plt.xlabel('x')
plt.ylabel('y')
```


But, True Model Order not Known

□ Suppose we guess the wrong model order?



$d=1$ "Underfitting"



$d=10$ "Overfitting"

How Can You Tell from Data?



- ❑ Is there a way to tell what is the correct model order to use?
- ❑ Must use the data. Do not have access to the true d ?
- ❑ What happens if we guess:
 - d too big?
 - d too small?

Using RSS on Training Data?

❑ Simple (but bad) idea:

- For each model order, d , find estimate $\hat{\beta}$
- Compute predicted values on training data

$$\hat{y}_i = \hat{\beta}^T x_i$$

- Compute RSS

$$RSS(d) = \sum_i (y_i - \hat{y}_i)^2$$


- Find d with lowest RSS

❑ This doesn't work

- $RSS(d)$ is always decreasing (Question: Why?)
- Minimizing $RSS(d)$ will pick d as large as possible
- Leads to overfitting


❑ What went wrong?

❑ How do we do better?



Outline

☐ Motivating Example: What polynomial degree should a model use?

 ☐ Bias and variance

☐ Bias and variance for linear models (Advanced)

☐ Cross-validation

☐ Feature selection

Model Class

□ Consider general estimation problem

- Given data (x_i, y_i) want to learn a functional relation: $y \approx \hat{y} = f(x)$

□ Model class: The set of possible estimates:

$$\hat{y} = f(x, \beta)$$

- Set is parametrized by β

□ Many possible examples:

- Linear model: $\hat{y} = \beta_0 + \beta_1 x_1 + \dots + \beta_k x_k$
- Polynomial model: $\hat{y} = \beta_0 + \beta_1 x + \dots + \beta_k x^k$
- Nonlinear: $\hat{y} = \beta_0 + \beta_1 e^{-\beta_2 x} + \beta_3 e^{-\beta_4 x}$
- ...

Model Class and True Function

□ Analysis set-up:

- Learning algorithm assumes a **model class**: $\hat{y} = f(x, \beta)$
- But, data has **true** relation: $y = f_0(x) + \epsilon$, $\epsilon \sim N(0, \sigma_\epsilon^2)$

□ Will quantify three key effects:

- Irreducible error
- Under-modeling
- Over-fitting

Output Mean Squared Error

□ To evaluate prediction error suppose we are given:

- A parameter estimate $\hat{\beta}$ (computed from the learning algorithm for a fixed training set)
- A test point x_{test}
- Test point is generally different from training samples.

□ Predicted value: $\hat{y} = f(x_{test}, \hat{\beta})$

□ Actual value: $y = f_0(x_{test}) + \epsilon$

□ Define output mean squared error given $\hat{\beta}$:

$$MSE_y(x_{test}, \hat{\beta}) := E[y - \hat{y}]^2$$

- Expectation is over noise ϵ on the test sample.

Irreducible Error

□ Rewrite output MSE:

$$MSE_y(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}}) := E[y - \hat{y}]^2 = E[f_0(\mathbf{x}_{test}) + \epsilon - f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}})]^2$$

□ Since noise on test sample is independent of $\hat{\boldsymbol{\beta}}$ and \mathbf{x}_{test} :

$$MSE_y(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}}) := [f_0(\mathbf{x}_{test}) - f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}})]^2 + E(\epsilon^2) = [f_0(\mathbf{x}_{test}) - f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}})]^2 + \sigma_\epsilon^2$$

□ Define **irreducible error**: σ_ϵ^2

- Lower bound on $MSE_y(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}}) \geq \sigma_\epsilon^2$
- Fundamental limit on ability to predict y
- Occurs since y is influenced by other factors than \mathbf{x}

Under-Modeling

□ **Definition:** A true function $f_0(x)$ is **in the model class** $\hat{y} = f(x, \beta)$ if:

$$f_0(x) = f(x, \beta_0) \text{ for all } x$$

for some parameter β_0 .

- β_0 called the **true parameter**

□ **Under-modeling:** When $f_0(x)$ is not in the model class

Sample Question

□ For each pair, state if the true function is in the model class or not

- That is, is there under-modeling or not?
- If true function is in the model class, state the true parameter

□ Examples:

- True function: $f_0(x) = 2 + 3x$ Model class: $f(x, \beta) = \beta_0 + \beta_1 x + \beta_2 x^2$
- True function: $f_0(x) = 2 + 3x + 4x^2$ Model class: $f(x, \beta) = \beta_0 + \beta_1 x$
- True function: $f_0(x) = \sin(2\pi(5)x + 7)$ Model class: $f(x, \beta) = \beta_0 \sin(2\pi(5)x) + \beta_1 \cos(2\pi(5)x)$
- True function: $f_0(x) = \sin(2\pi(8)x + 7)$ Model class: $f(x, \beta) = \beta_0 \sin(2\pi(5)x) + \beta_1 \cos(2\pi(5)x)$

□ Solutions in class

Under-Modeling and Irreducible Error

□ Suppose that:

- There is no under-modeling: $f_0(\mathbf{x}) = f(\mathbf{x}, \boldsymbol{\beta}_0)$ for some “true” parameter $\boldsymbol{\beta}_0$; and
- Estimator selects the true parameter $\hat{\boldsymbol{\beta}} = \boldsymbol{\beta}_0$

□ Then, output error is:

$$MSE_y(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}}) := [f_0(\mathbf{x}_{test}) - f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}})]^2 + \sigma_\epsilon^2 = \sigma_\epsilon^2$$

□ **Conclusion:** If there is no undermodeling and we can estimate the true parameter:

- We can get output error = irreducible error
- We can achieve the same error as if we knew the true function $f_0(\mathbf{x})$

Bias of an Estimator

- Suppose training data (\mathbf{x}_i, y_i) is generated as follows:
 - Fix data input points $\mathbf{x}_i, i = 1, \dots, N$ (Treat as non-random)
 - Generate data output points $y_i = f_0(\mathbf{x}_i) + \epsilon_i$ with random i.i.d. noise ϵ_i with some distribution
- Then estimate $\hat{\boldsymbol{\beta}}$ is a random vector
 - Depends on the noise ϵ_i in the training data
- **Definition:** The **bias** at a test point \mathbf{x}_{test} is:

$$\text{Bias}(\mathbf{x}_{test}) := f_0(\mathbf{x}_{test}) - E[f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}})]$$

- Measures the difference between:
 - True function $f_0(\mathbf{x}_{test})$
 - Expected value of estimated $f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}})$, averaged over the noise in the training data

Bias: Noise-Free Case

□ Suppose true relation has no noise: $y = f_0(\mathbf{x})$

- Will handle noise later

□ Get training data: $(\mathbf{x}_i, y_i), i = 1, \dots, n$

□ Fit model parameter from least-squares:

$$\hat{\boldsymbol{\beta}} = \arg \min_{\boldsymbol{\beta}} \sum_{i=1}^n (y_i - f(\mathbf{x}_i, \boldsymbol{\beta}))^2 = \arg \min_{\boldsymbol{\beta}} \sum_{i=1}^n (f_0(\mathbf{x}_i) - f(\mathbf{x}_i, \boldsymbol{\beta}))^2$$

- Minimizing training error finds best least squares fit of the true functions in the model class

□ **Conclusions:** If

- There is no under-modeling: $f_0(\mathbf{x}_i) = f(\mathbf{x}_i, \boldsymbol{\beta}_0)$ for some true parameter $\boldsymbol{\beta}_0$
- Minimization for $\hat{\boldsymbol{\beta}}$ is unique

Then $\hat{\boldsymbol{\beta}} = \boldsymbol{\beta}_0$ and $\text{Bias}(\mathbf{x}_{test}) = 0$ for all \mathbf{x}_{test} :

- No bias when there is no under-modeling and no noise

□ Will show later that for linear models, there is no bias even when there is no noise

Bias Visualized


- Polynomial example

 - $d_{true} = 3$

- No noise in data

- No bias when
 $d \geq d_{true}$

Bias(x)



MSE of an Estimator

- Data model: $y = f_0(\mathbf{x}) + \epsilon, \epsilon \sim N(0, \sigma_\epsilon^2)$
- Get training data: $(\mathbf{x}_i, y_i), i = 1, \dots, n$
- Fit parameter $\hat{\boldsymbol{\beta}}$ from data (e.g. via least squares)
 - $\hat{\boldsymbol{\beta}}$ will be random. Depends on particular noise realization for the selected training samples.
- Take a new test point \mathbf{x}_{test}
- Define two mean square errors:
 - Output MSE: $\text{MSE}_y(\mathbf{x}_{test}) := E[y - f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}})]$: Error on the predicted value
 - Function MSE: $\text{MSE}_f(\mathbf{x}_{test}) := E[f_0(\mathbf{x}_{test}) - f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}})]$: Error on the underlying function
- Expectation is over both:
 - Noise in the training data : $y_i = f_0(\mathbf{x}_i) + \epsilon_i$
 - Noise on the test sample: $y = f_0(\mathbf{x}_{test}) + \epsilon$

MSE and the Irreducible Error

□ From previous slide:

- Output MSE: $MSE_y(\mathbf{x}_{test}) := E[y - f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}})]$: Error on the predicted value
- Function MSE: $MSE_f(\mathbf{x}_{test}) := E[f_0(\mathbf{x}_{test}) - f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}})]$: Error on the underlying function

□ **Theorem:** $MSE_y(\mathbf{x}_{test}) = MSE_f(\mathbf{x}_{test}) + \epsilon^2$

- Recall ϵ^2 is the irreducible error

□ **Proof:** Similar to before:

- We know $y = f_0(\mathbf{x}_{test}) + \epsilon$
- $MSE_y(\mathbf{x}_{test}) = E[y - f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}})]^2 = E[f_0(\mathbf{x}_{test}) + \epsilon - f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}})]^2$
- But, ϵ is independent of $f_0(\mathbf{x}_{test})$ and $f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}})$
- Therefore $MSE_y(\mathbf{x}_{test}) = E[f_0(\mathbf{x}_{test}) - f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}})]^2 + E(\epsilon^2) = MSE_f(\mathbf{x}_{test}) + \sigma_\epsilon^2$

Bias and Variance

□ We will show that function MSE can be related to two key quantities

□ **Bias:** $Bias(\mathbf{x}_{test}) := f_0(\mathbf{x}_{test}) - E[f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}})]$

- How much the average value of the estimate differs from the true function

□ **Variance:** $Var(\mathbf{x}_{test}) := E \left[f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}}) - E[f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}})] \right]^2$

- How much the estimate varies around its average

□ Bias and variance are (conceptually) measured as follows:


- Get many independent training data sets, each with same size N and input values \mathbf{x}_i
- Each dataset has different output values y_i because of independent noise in the training data
- Obtain $\hat{\boldsymbol{\beta}}$ for each training data set
- Bias and variances are computed over the different sets

□ Of course, in reality, we have only one training dataset

□ But, bias and variance are used to study theoretical averages over different experiments

Bias and Variance Illustrated

- Polynomial ex
- Mean and std dev of estimated functions
- 100 trials
- Solid line: mean estimate among all trials
- Error bars: 1 STD



Bias-Variance Formula

□ Recall definitions:

- **Function MSE:** $MSE_f(\mathbf{x}_{test}) := E[f_0(\mathbf{x}_{test}) - f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}})]:$
- **Bias:** $Bias(\mathbf{x}_{test}) := f_0(\mathbf{x}_{test}) - E[f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}})]$
- **Variance:** $Var(\mathbf{x}_{test}) := E[f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}}) - E[f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}})]]^2$

□ Bias-Variance formula : $MSE_f(\mathbf{x}_{test}) = Bias(\mathbf{x}_{test})^2 + Var(\mathbf{x}_{test})$

- Will be proved below

□ Bias-Variance tradeoff:


□ Bias due to under-modeling

- Reduced with high model order

□ Variance is due to noise in training data and number of parameters to estimate

- Increases with higher model order

Bias-Variance Tradeoff



□ Bias:

- Due to under-modeling
- Reduced with high model order

□ Variance:

- Increases with noise in training data
- Increase with high model order

□ Optimal model order depends on:

- Amount of samples available
- Underlying complexity of the relation

Bias-Variance Formula Proof


□ Define $\bar{f}(\mathbf{x}_{test}) = E[f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}})]$ = average value of estimated function

□ $MSE_f(\mathbf{x}_{test}) = E[f_0(\mathbf{x}_{test}) - f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}})]^2 = E[f_0(\mathbf{x}_{test}) - \bar{f}(\mathbf{x}_{test}) + \bar{f}(\mathbf{x}_{test}) - f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}})]^2$

□ Three components: $MSE_f(\mathbf{x}_{test}) = M_1 + M_2 - 2M_3$

- $M_1 = E[f_0(\mathbf{x}_{test}) - \bar{f}(\mathbf{x}_{test})]^2 = [f_0(\mathbf{x}_{test}) - \bar{f}(\mathbf{x}_{test})]^2 = Bias(\mathbf{x}_{test})$
- $M_2 = E[f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}}) - \bar{f}(\mathbf{x}_{test})]^2 = Var(\mathbf{x}_{test})$
- $M_3 = E[(f_0(\mathbf{x}_{test}) - \bar{f}(\mathbf{x}_{test}))(f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}}) - \bar{f}(\mathbf{x}_{test}))]$
 $= (f_0(\mathbf{x}_{test}) - \bar{f}(\mathbf{x}_{test}))E[f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}}) - \bar{f}(\mathbf{x}_{test})]$
 $= (f_0(\mathbf{x}_{test}) - \bar{f}(\mathbf{x}_{test}))(f(\mathbf{x}_{test}) - \bar{f}(\mathbf{x}_{test})) = 0$

Outline

- ☐ Motivating Example: What polynomial degree should a model use?
- ☐ Bias and variance
-  ☐ Bias and variance for linear models (Advanced)
- ☐ Cross-validation
- ☐ Feature selection

This Section is Advanced

- ❑ This section requires more advanced probability and linear algebra
- ❑ Means and variances of random vectors
- ❑ Undergraduates: Skip to final slide for final conclusions
- ❑ Graduate students: We will cover this
 - You should review your multi-variable probability and linear algebra

Linear Models

- Consider linear model in general transformed feature space:

$$\hat{y} = f(x, \beta) = \phi(x)^T \beta = \beta_1 \phi_1(x) + \cdots + \beta_p \phi_p(x)$$

- See previous lecture

- Assume true data relation is: $y = f_0(x) + \epsilon$, $E(\epsilon) = 0$, $E(\epsilon^2) = \sigma^2$

- When there is no under-modeling: $f_0(x) = f(x, \beta^0) = \phi(x)^T \beta^0$

- $\beta^0 = (\beta_0^0, \dots, \beta_k^0)$ True parameter

- Get data (x_i, y_i) , $i = 1, \dots, N$

- Least squares fit $\hat{\beta} = (A^T A)^{-1} A^T y$

$$A = \begin{bmatrix} \phi_1(x_1) & \cdots & \phi_p(x_1) \\ \vdots & \vdots & \vdots \\ \phi_1(x_N) & \cdots & \phi_p(x_N) \end{bmatrix}$$

Minimum Number of Samples

- ❑ LS estimate requires $A^T A$ is invertible.
- ❑ Linear algebra fact: Since $A \in R^{N \times p}$, we need $\text{Rank}(A) \geq p$
 - Otherwise solution is not unique
- ❑ Since $\text{Rank}(A) \leq \min(N, p)$ we need $N \geq p$.
- ❑ Recall:
 - N = number of data samples
 - p = number of parameters
- ❑ **Conclusion:** Number of samples \geq number of parameters
- ❑ This places a basic limit on the model complexity that you can use

Random Vectors Review

- To analyze bias and variance in linear models, we need to review random vectors
- **Random vectors:** $\mathbf{x} = (x_1, \dots, x_d)^T$: Each component x_j is a random variable
- **Mean:** The vector of means of the components

$$\boldsymbol{\mu} = E\mathbf{x} = (Ex_1, \dots, Ex_d)^T = (\mu_1, \dots, \mu_d)^T$$

- **Covariance components:** $\text{Cov}(x_i, x_j) = E[(x_i - \mu_i)(x_j - \mu_j)]$

- **Variance matrix ($d \times d$):**

$$\text{Var}(\mathbf{x}) := E[(\mathbf{x} - \boldsymbol{\mu})(\mathbf{x} - \boldsymbol{\mu})^T] = \begin{bmatrix} \text{Cov}(x_1, x_1) & \cdots & \text{Cov}(x_1, x_d) \\ \vdots & \ddots & \vdots \\ \text{Cov}(x_d, x_1) & \cdots & \text{Cov}(x_d, x_d) \end{bmatrix}$$

Linear Transforms of Random Vectors

- A linear transform is a map: $y = Ax + b$
- $A \in R^{M \times N}$ maps input $x \in R^N$ to $Ax \in R^M$
- Mean and variance matrix under linear map given by
 - Mean: $E(y) = AE(x) + b$
 - Variance: $Var(y) = AVar(x)A^T$

Bias With No Under-Modeling

- Suppose that there is no undermodeling: $f_0(x) = \phi(x)^T \beta^0$
- Then each training sample output is: $y_i = \phi(x_i)^T \beta^0 + \epsilon_i$
- Hence: true data vector $y = A\beta^0 + \epsilon$
- Parameter estimate is:
$$\hat{\beta} = (A^T A)^{-1} A^T y = (A^T A)^{-1} A^T (A\beta^0 + \epsilon) = \beta^0 + (A^T A)^{-1} A^T \epsilon$$
- Since $E\epsilon = 0$, $E\hat{\beta} = \beta^0$. Average of parameter estimate matches true parameter
- $Ef(x_{test}, \hat{\beta}) = \phi(x_{test})^T E\hat{\beta} = \phi(x_{test})^T \beta^0 = f_0(x_{test})$
- Therefore $Bias(x_{test}) := f_0(x_{test}) - Ef(x_{test}, \hat{\beta}) = 0$
- **Conclusion:** When the model is linear and there is no under-modeling, there is no bias

Variance of the Parameters in Linear Models

- Since ϵ_i are independent for different samples with $E\epsilon_i = 0$, $E\epsilon_i^2 = \sigma^2$

$$\text{Cov}(\epsilon_i, \epsilon_j) = \begin{cases} 0 & i \neq j \\ \sigma^2 & i = j \end{cases}$$

- Therefore variance matrix is: $\text{Var}(\epsilon) = \sigma^2 I$

- From last slide: $\hat{\beta} = \beta^0 + (A^T A)^{-1} A^T \epsilon$.

- Applying variance formula of a linear transformation of ϵ

$$\begin{aligned} E \left((\hat{\beta} - \beta^0)(\hat{\beta} - \beta^0)^T \right) &= (A^T A)^{-1} A^T \text{Var}(\epsilon) A (A^T A)^{-1} \\ &= \sigma^2 (A^T A)^{-1} A^T A (A^T A)^{-1} = \sigma^2 (A^T A)^{-1} \end{aligned}$$

Variance in Linear Estimate

- To compute variance use trick: Suppose \mathbf{a} and \mathbf{z} are vectors, \mathbf{a} is non-random, \mathbf{z} is random:

$$E|\mathbf{a}^T \mathbf{z}|^2 = E(\mathbf{a}^T \mathbf{z} \mathbf{z}^T \mathbf{a}) = \mathbf{a}^T E(\mathbf{z} \mathbf{z}^T) \mathbf{a}$$

- From earlier: $E f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}}) = \phi(\mathbf{x}_{test})^T E \hat{\boldsymbol{\beta}} = \phi(\mathbf{x}_{test})^T \boldsymbol{\beta}^0$

- Therefore variance of linear model:

$$\begin{aligned} \text{Var}(\mathbf{x}_{test}) &= E[f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}}) - E f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}})]^2 = E[\phi(\mathbf{x}_{test})^T (\hat{\boldsymbol{\beta}} - \boldsymbol{\beta}^0)]^2 \\ &= \phi(\mathbf{x}_{test})^T E[(\hat{\boldsymbol{\beta}} - \boldsymbol{\beta}^0)(\hat{\boldsymbol{\beta}} - \boldsymbol{\beta}^0)^T] \phi(\mathbf{x}_{test}) \\ &= \sigma^2 \phi(\mathbf{x}_{test})^T (\mathbf{A}^T \mathbf{A})^{-1} \phi(\mathbf{x}_{test}) \end{aligned}$$

- Above calculation is for the case of no under-modeling
- But, similar calculation shows variance expression is the same when there is under-modeling

Case with Equal Test & Training Distributions

□ Suppose that test point is distributed identically to training data

- Training data inputs $\mathbf{x}_i, i = 1, \dots, N$
- $\mathbf{x}_{test} = \mathbf{x}_i$ with probability $\frac{1}{N}$

□ Since rows of A are $\phi(\mathbf{x}_i)^T$: $A^T A = \sum_{i=1}^N \phi(\mathbf{x}_i) \phi(\mathbf{x}_i)^T$

□ Now use trick: For random vectors \mathbf{u}, \mathbf{v} : $E(\mathbf{u}^T \mathbf{v}) = \text{Tr } E(\mathbf{v} \mathbf{u}^T)$

- $\text{Tr}(A) = \sum_i A_{ii}$ = sum of diagonals

□ Therefore, variance averaged over \mathbf{x}_{test} is:

$$E \text{Var}(\mathbf{x}_{test}) = \sigma^2 E[\phi(\mathbf{x}_{test})^T (A^T A)^{-1} \phi(\mathbf{x}_{test})] = \sigma^2 \text{Tr}\{E[\phi(\mathbf{x}_{test}) \phi(\mathbf{x}_{test})^T] (A^T A)^{-1}\}$$

$$= \frac{\sigma^2}{N} \text{Tr} \left\{ \sum_i \phi(\mathbf{x}_i) \phi(\mathbf{x}_i)^T (A^T A)^{-1} \right\} = \frac{\sigma^2}{N} \text{Tr}\{(A^T A)(A^T A)^{-1}\} = \frac{\sigma^2}{N} \text{Tr}\{I_p\} = \frac{\sigma^2 p}{N}$$

Case with Equal Test & Training Distributions

- Assumption on previous slide: Test point \mathbf{x}_{test} is randomly selected from training data
- Then, average variance is given by

$$E \text{Var}(\mathbf{x}_{test}) = \frac{\sigma^2 p}{N}$$

- Increases with number of parameters p
 - Shows that increasing model complexity increases variance error
- Decreases with number of samples N
- What if test data point is distributed differently from training data?
 - Then variance may be much larger $\frac{\sigma^2 p}{N}$
 - If test data is not like training data, we are extending model to regions not seen in training data
 - Often leads to high error

Summary of Results for Linear Models

□ Suppose model is linear with N = num samples, p = num parameters

□ Result 1: When $N < p$, linear estimate is not unique

- Need at least as many samples as parameters

□ Now assume that $N \geq p$ and parameter estimate is unique

□ Result 2: When there is no under-modeling, estimate is unbiased


$$E[f(x_{test}, \hat{\beta})] = f_0(x_{test}).$$

□ Result 3: If test point drawn from same distribution as training data:

$$Var = \frac{p}{n} \sigma_{\epsilon}^2$$

- Variance increases linearly with number of parameters and inversely with number of samples

Outline

- ☐ Motivating Example: What polynomial degree should a model use?
- ☐ Bias and variance
- ☐ Bias and variance for linear models (Advanced)
-  ☐ Cross-validation
- ☐ Feature selection

Cross Validation

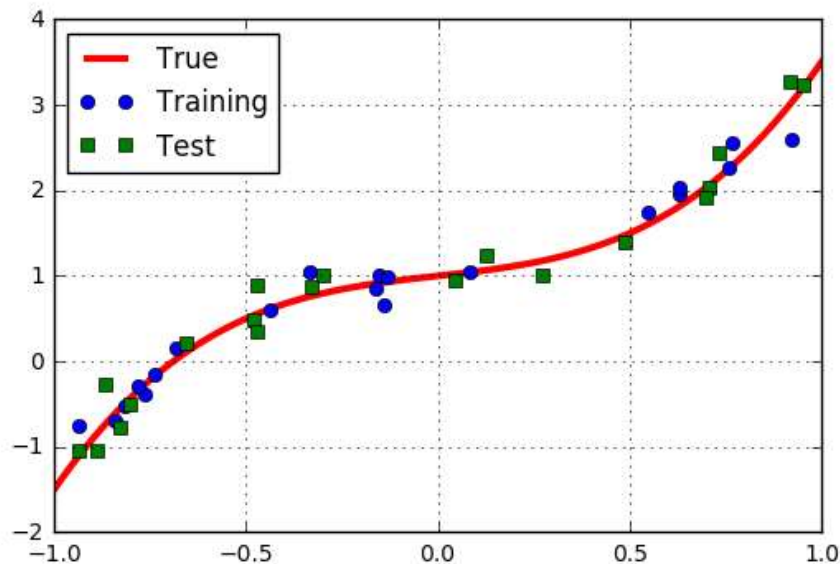
- ❑ Concept: Need to fit on test data independent of training data
- ❑ Divide data into two sets:
 - N_{train} training samples, N_{test} test samples
- ❑ For each model order, p , learn parameters $\hat{\beta}$ from training samples
- ❑ Measure RSS on test samples.

$$RSS_{test}(p) = \sum_{i \in \text{test}} (\hat{y}_i - y_i)^2$$

- ❑ Select model order p that minimizes $RSS_{test}(p)$

Polynomial Example: Training Test Split

Example: Split data into 20 samples for training, 20 for test




```
# Number of samples for training and test
ntr = nsamp // 2
nts = nsamp - ntr

# Training
xtr = xdat[:ntr]
ytr = ydat[:ntr]

# Test
xts = xdat[ntr:]
yts = ydat[ntr:]
```

Finding the Model Order

Estimated optimal model order = 3



RSS test minimized at $d = 3$
RSS training always decreases

```
dtest = np.array(range(0,10))
RSS_test = []
RSS_train = []
for d in dtest:

    # Fit data
    beta_hat = poly.polyfit(xtr,ytr,d)

    # Measure RSS on training data
    # This is not necessary, but we do it just to show the training error
    yhat = poly.polyval(xtr,beta_hat)
    RSSd = np.mean((yhat-ytr)**2)
    RSS_train.append(RSSd)

    # Measure RSS on test data
    yhat = poly.polyval(xts,beta_hat)
    RSSd = np.mean((yhat-yts)**2)
    RSS_test.append(RSSd)

plt.plot(dtest,RSS_train,'bo-')
plt.plot(dtest,RSS_test,'go-')
plt.xlabel('Model order')
plt.ylabel('RSS')
plt.grid()
plt.ylim(0,1)
plt.legend(['Training','Test'],loc='upper right')
```

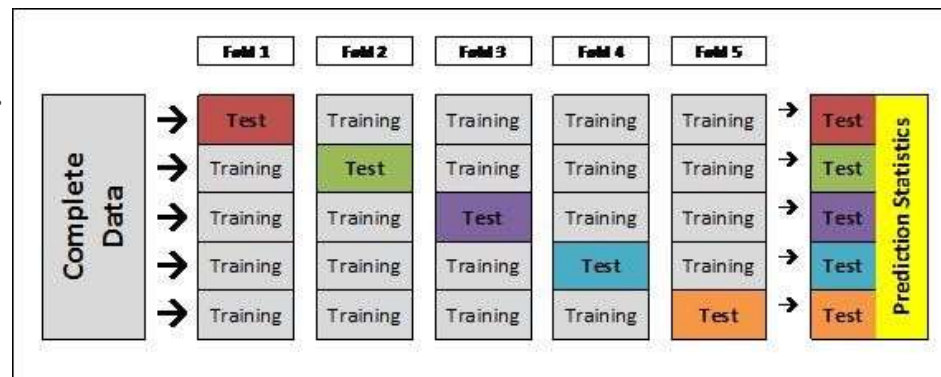
Problems with Simple Train/Test Split

- ❑ Test error could vary significantly depending on samples selected
- ❑ Only use limited number of samples for training
- ❑ Problems particularly bad for data with limited number of samples

K-Fold Cross Validation

□ K-fold cross validation

- Divide data into K parts
- Use $K - 1$ parts for training. Use remaining for test.
- Average over the K test choices
- More accurate, but requires K fits of parameters
- Typical choice: $K=5$ or 10
- Average MSE over K folds estimates the total MSE
- $(= \text{Bias}^2 + \text{Variance} + \text{irreducible error})$



□ Leave one out cross validation (LOOCV)

- Take $K = N$ so one sample is left out.
- Most accurate, but requires N model fittings
- Necessary when N is small.

From

<http://blog.goldenhelix.com/goldenadmin/cross-validation-for-genomic-prediction-in-svs/>

Polynomial Example

❑ Use sklearn Kfold object

❑ Loop

- Outer loop: Over K folds
- Inner loop: Over D model orders
- Measure test error in each fold and order
- Averaging test errors from K folds for each model order
- Find the model order with the minimal average test errors
- Can be time-consuming

```
# Create a k-fold object
nfold = 20
kf = sklearn.model_selection.KFold(n_splits=nfold, shuffle=True)

# Model orders to be tested
dtest = np.arange(0,10)
nd = len(dtest)

# Loop over the folds
RSSs = np.zeros((nd,nfold))
for isplit, Ind in enumerate(kf.split(xdat)):

    # Get the training data in the split
    Itr, Its = Ind
    xtr = xdat[Itr]
    ytr = ydat[Itr]
    xts = xdat[Its]
    yts = ydat[Its]

    for it, d in enumerate(dtest):


        # Fit data on training data
        beta_hat = poly.polyfit(xtr,ytr,d)

        # Measure RSS on test data
        yhat = poly.polyval(xts,beta_hat)
        RSSs[it,isplit] = np.mean((yhat-yts)**2)
```


Polynomial Example CV Results


- For each model order d
 - Compute mean test RSS over K folds
 - Compute standard error (SE) of test RSS
 - $SE = \text{STD of mean RSS} = \text{RSS std} / \sqrt{K - 1}$
 - (expectation over different realizations of data in each fold)
- Simple model selection
 - Select d with lowest mean test RSS
- For this example
 - Estimate model order = 3

```
RSS_mean = np.mean(RSSsts,axis=1)
RSS_std = np.std(RSSsts,axis=1) / np.sqrt(nfold-1)
plt.errorbar(dtest, RSS_mean, yerr=RSS_std, fmt='-')
plt.ylim(0,1)
plt.xlabel('Model order')
plt.ylabel('Test RSS')
plt.grid()
```




One Standard Error Rule

- Previous slide: Select d to minimize $\text{mse_mean}[d]$
- Problem: Often over-predicts model order
- One standard deviation rule
 - Use simplest model within one SE of minimum
- Detailed procedure:
 - Find d_0 to minimize $\text{mse_mean}[d]$
 - Set $\text{mse_tgt} = \text{mse_mean}[d_0] + \text{mse_std}[d_0]$
 - Find minimal d_{opt} s.t. $\text{mse_mean}[d_{\text{opt}}] \leq \text{mse_tgt}$



Outline

- ❑ Motivating Example: What polynomial degree should a model use?
- ❑ Bias and variance
- ❑ Bias and variance for linear models (Advanced)
- ❑ Cross-validation
- ❑ Feature selection

Feature Selection as Model Selection

- ❑ In many problems, we have a large number of features x_1, \dots, x_p
 - But only a small number of features are likely relevant
- ❑ Fit of linear model on all features may be impossible
 - Too many parameters
- ❑ **Feature selection** problem: Select a subset $S \subseteq \{1, \dots, p\}$ of $k \ll p$ features
- ❑ Can fit a linear model for each subset S using linear regression
- ❑ Model order is $k = \text{number of features to use}$
- ❑ Higher feature number leads to low bias but higher variance!

Exhaustive Search for Feature Selection

- Given potential features x_1, \dots, x_p
- **Problem:** Select subset S of features with $|S| \leq k$
- Simple **exhaustive search:**
 - Divide data into training test
 - For each subset $|S| \leq k$:
 - Fit a model with features in S
 - Measure test error $RSS(S)$
 - Select subset S with minimum test $RSS(S)$

Feature Selection based on Correlation

- ❑ Exhaustive search may be infeasible when the raw feature dimension is large!
 - There are $\binom{p}{k}$ subsets with size k
- ❑ Suboptimal approach:
 - For each candidate feature order $k \leq p$, choose k features with the highest correlation coefficients with the target
 - Use cross validation to determine the RSS mean and variance for this subset
 - Select the feature subset with minimal RSS mean or using the one standard error rule.
- ❑ Is using correlation with target a good idea?
 - Two features that are correlated could both be highly correlated with the target
 - Hence, they may provide redundant information and ideally only one of them should be used.

Greedy Feature Selection

❑ Forward-Stepwise Selection

- Select one feature from all features that provides the lowest RSS with cross validation
- Select one new feature from all remaining features, so that previously chosen features plus the new feature provides the lowest RSS
- Repeat until the maximum feature number is reached, or when the RSS starts to increase

❑ Backward-Stepwise

- First use all features and find the RSS (using cross validation)
- Remove one feature and find the new RSS. Go through all possible features to remove.
- Find the one that leads to the least RSS increase. Remove this feature.
- Repeat the above, remove one from the remaining features, to find the next most important feature.

❑ Except exhaustive search, can all lead to suboptimal solution

❑ We will discuss the method of LASSO in the next lecture for feature selection.

Comparison of Feature selection Methods

Figure from [Hastie2008]: Hastie, Tibshirani, Friedman, The elements of statistical learning.

For more on this subject, see Sec. 3.3


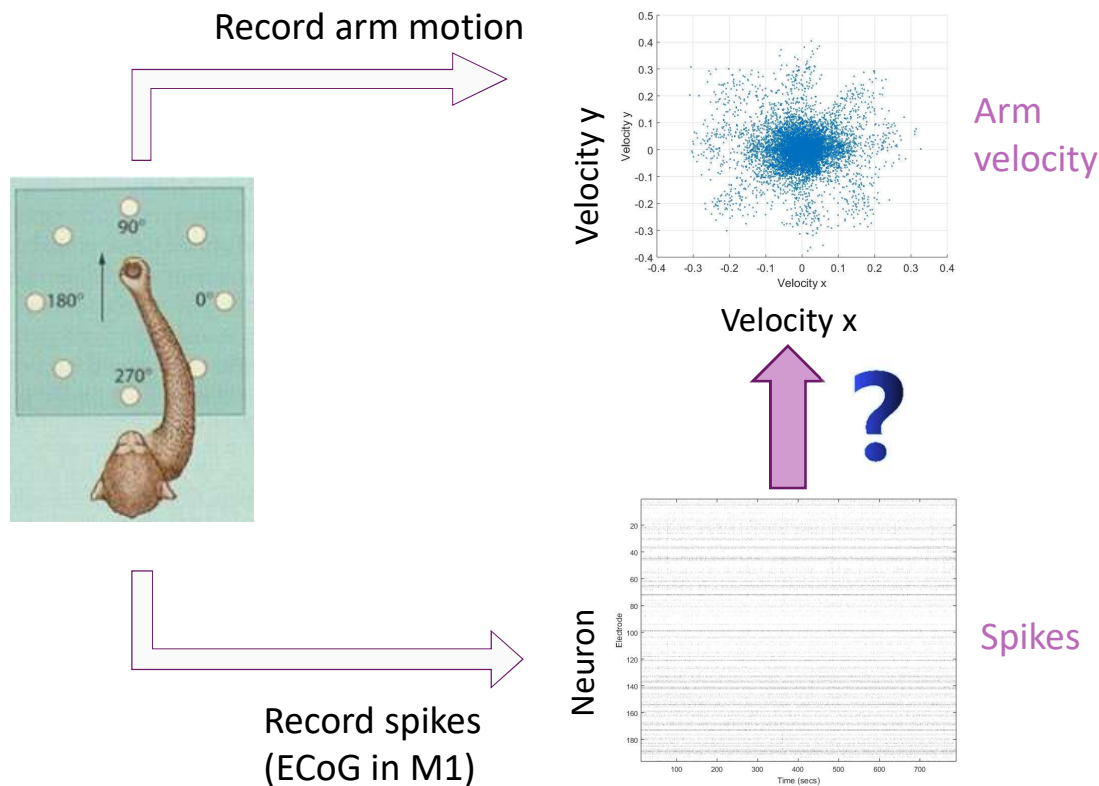


FIGURE 3.6. Comparison of four subset-selection techniques on a simulated linear regression problem $Y = X^T \beta + \varepsilon$. There are $N = 300$ observations on $p = 31$ standard Gaussian variables, with pairwise correlations all equal to 0.85. For 10 of the variables, the coefficients are drawn at random from a $N(0, 0.4)$ distribution; the rest are zero. The noise $\varepsilon \sim N(0, 6.25)$, resulting in a signal-to-noise ratio of 0.64. Results are averaged over 50 simulations. Shown is the mean-squared error of the estimated coefficient $\hat{\beta}(k)$ at each step from the true β .

Lab: Neural ECoG Data



- ❑ Read a monkey's brain!
- ❑ Predict hand motion from electrode measurements (number of spikes in each neuron)
- ❑ Feature selection
 - Which neuron signals predict arm motion?
 - Use correlation to select candidate feature subsets
- ❑ Data from <https://crcns.org/>
 - Open source website on neural data
 - Great for projects