Lecture 13
Decision Trees and Random Forest

EE-UY 4563/EL-GY 9123: INTRODUCTION TO MACHINE LEARNING

PROF. YAO WANG

Outline

❑Decision tree as constrained space partition

❑Regression tree design

❑Decision tree pruning

❑Classification tree design

❑Bagging

❑Random Forest

❑Feature ranking from random forest

Decision tree as constrained space partitioning

- Each region is regressed/classified to the same value
- The partition can be specified sequentially by splitting the range of one feature at a time.
- The splitting rule can be described by a tree.
  - Each leaf node = One region
  - Size of tree $|T|$ = number of leaf nodes
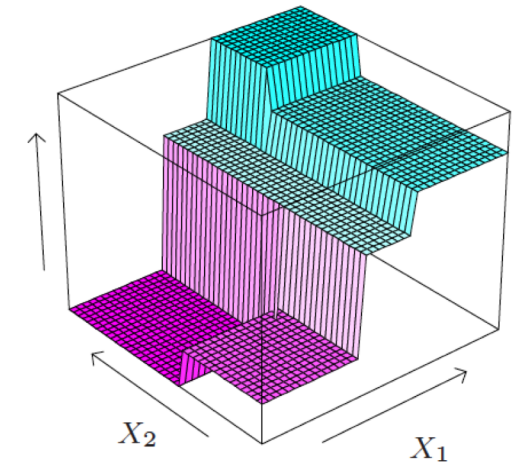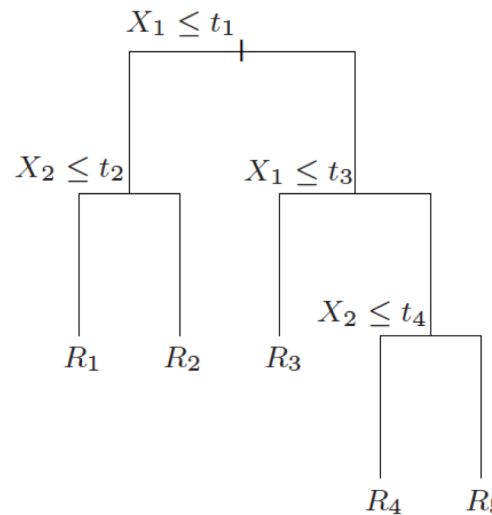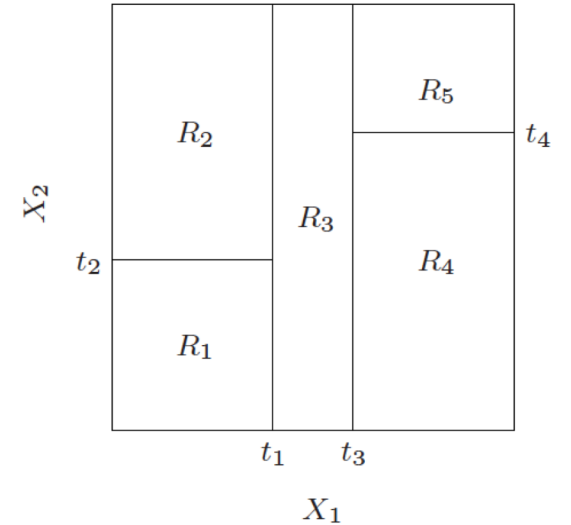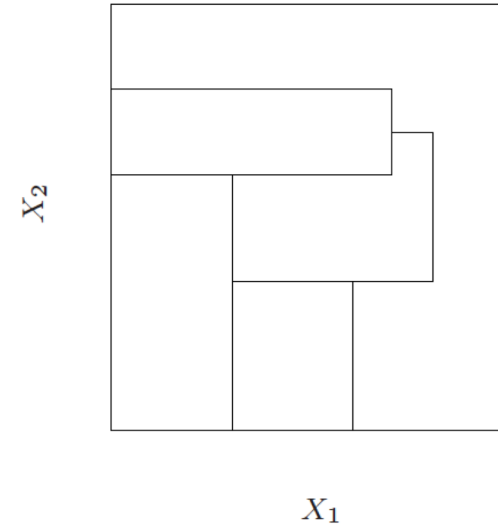- The partition is constrained: only rectangles in the 2D case.

Fig. 9.2 in ESL

How to build a decision tree? (Regression Case)

❑ Goal: minimize RSS

$$L = \sum_{m=1}^{|T|} \sum_{x_n \in R_m} (y_n - \bar{y}_m)^2$$

❑ Greedy algorithm:

❑ Start with a single region (entire space) and iterate:

❑ For each region R_m, select a feature x_j, and a splitting threshold s, such that splitting R_k with the criterion $x_j < s$ produces the largest decrease in RSS in R_m

  ◦ Exhaustive search: for each x_j, try all possible s in the current range of x_j in R_m

❑ Stop splitting a region if it contains <= N_{min} samples

R_1 R_2
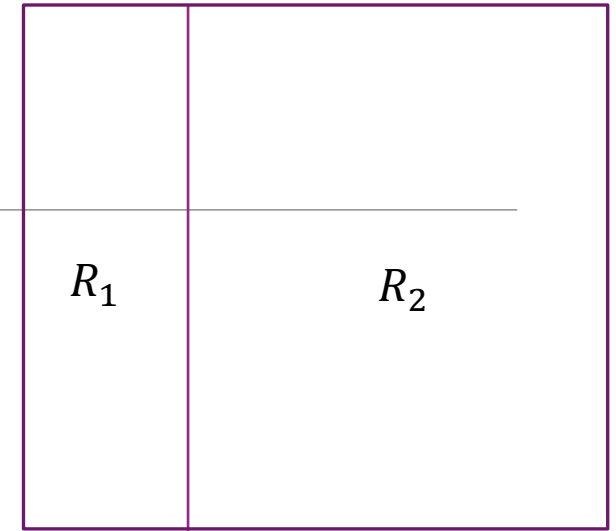
All possible splits of R_1:

R_2
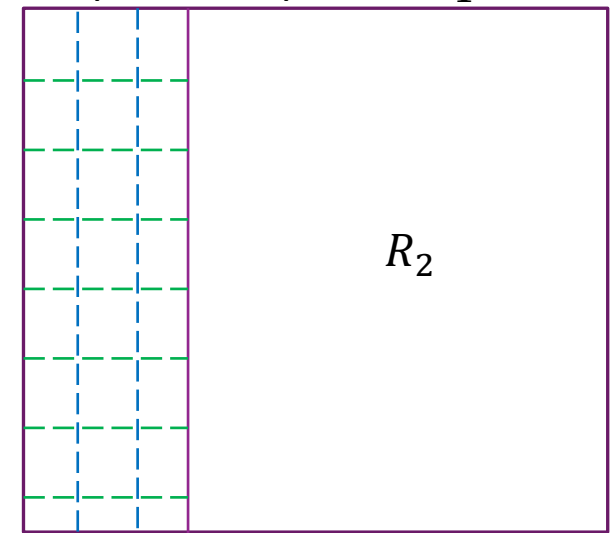
Overfitting

❑ Decision tree is very prone to overfitting

❑ Can exactly represent any function defined by the training set by having as many regions (or leaf nodes) as needed (Fully grown tree)

❑ How to control overfitting?
- Find optimal subtree (with a certain constraint on the minimum number of samples in the leaf nodes or maximum depth) by cross validation: too many possibilities
- Stop growing once RSS stop decreasing by a threshold with any new cut:
  - Not good because we use greedy search. It is possible to find a good cut after a bad one.
- Better idea: grow a full tree first, then prune the tree.

Weakest link pruning

❑ Starting with with the initial full tree T0 , merge two adjacent leave nodes (daughter nodes) to a single leaf node (mother). Select which nodes to merge by minimizing error increase. This produces a tree with one less region (or node)

❑ Repeat to merge another two nodes, until the minimum size tree is reached (e.g. a stump with 2 nodes)

❑ Generate a sequence of trees

T0, T1, T2, T3, …

❑ Which one to choose?
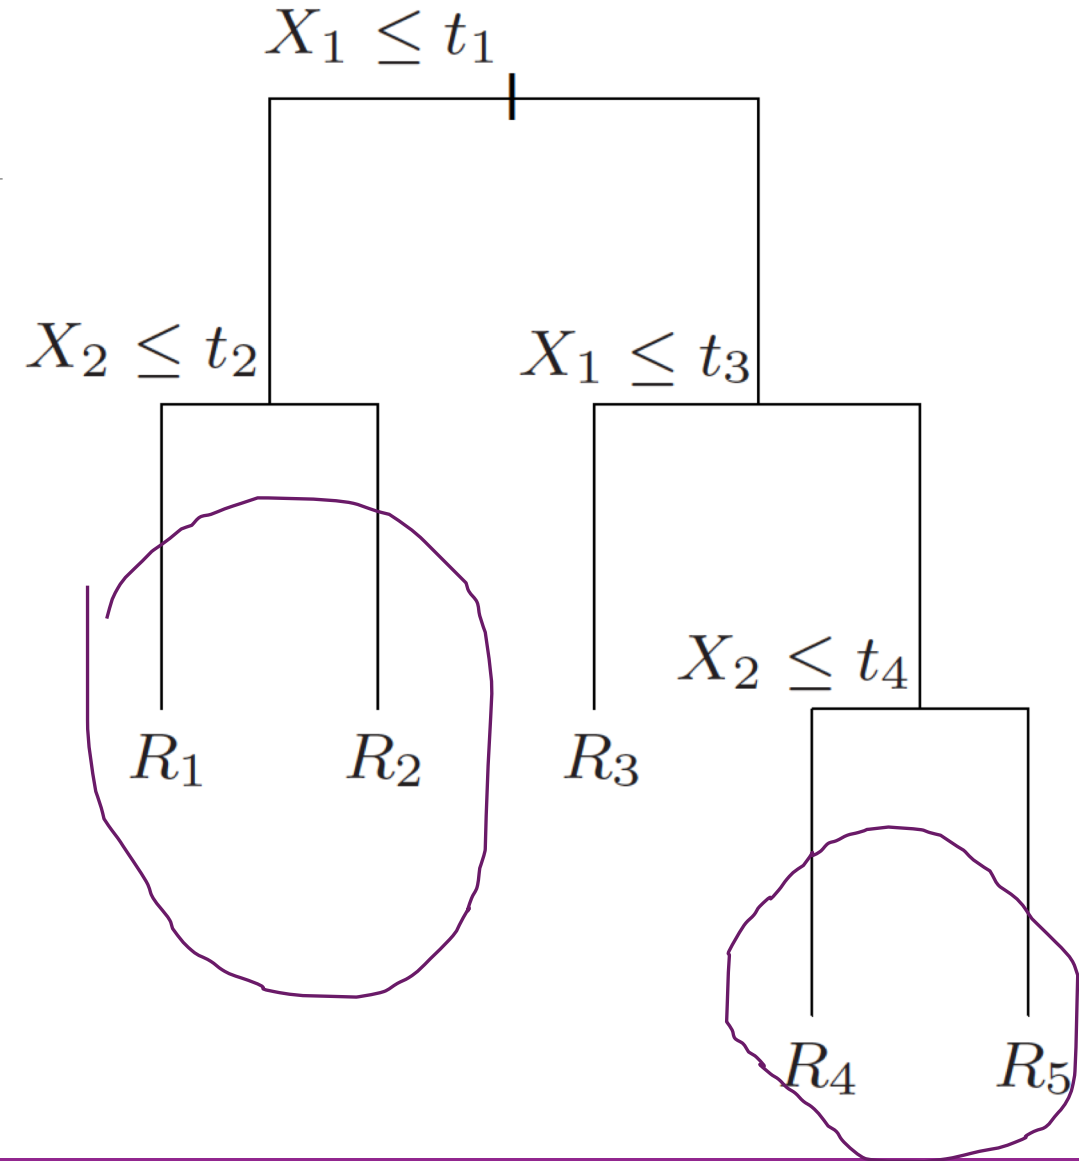
Cost complexity pruning

❑ Minimize a complexity regularized loss, over all possible trees T0, T1, T2, …

$$L(T, \alpha) = \sum_{m=1}^{|T|} \sum_{x_n \in R_m} (y_n - \bar{y}_m)^2 + \alpha|T|$$

Number of leaf nodes (regions)

- ◦ α=0: Full tree, $\alpha = \infty$: minimum sized tree

❑ How to choose α? Cross validation!

- ◦ For each α
  - ◦ For each validation fold:
    - ◦ build a sequence of trees using the training set, and finding the RSS on the testing set for each candidate tree. Find a tree that minimizes $L(T, \alpha)$
  - ◦ Find average $L(T, \alpha)$ over all validation folds

❑ When dataset is very large, can just pick one tree that has minimal RSS for the testset.

Example: Predicting baseball player salaries

From http://web.stanford.edu/class/stats202/content/lec19.pdf
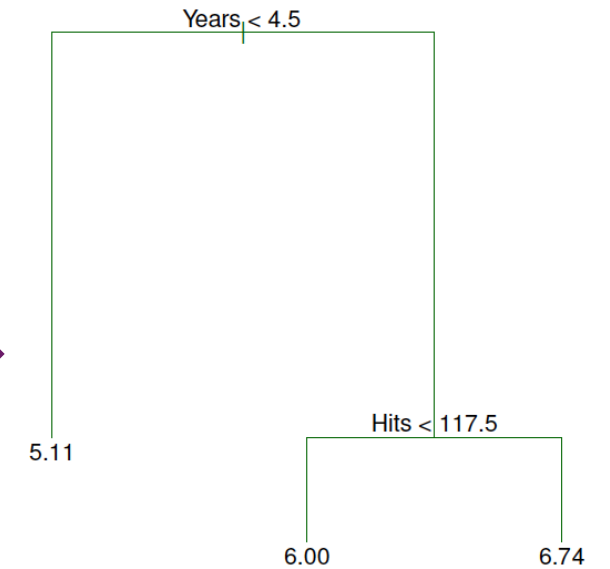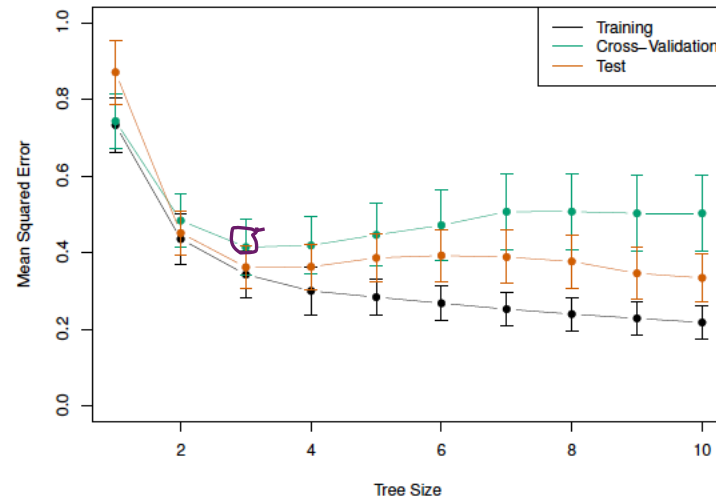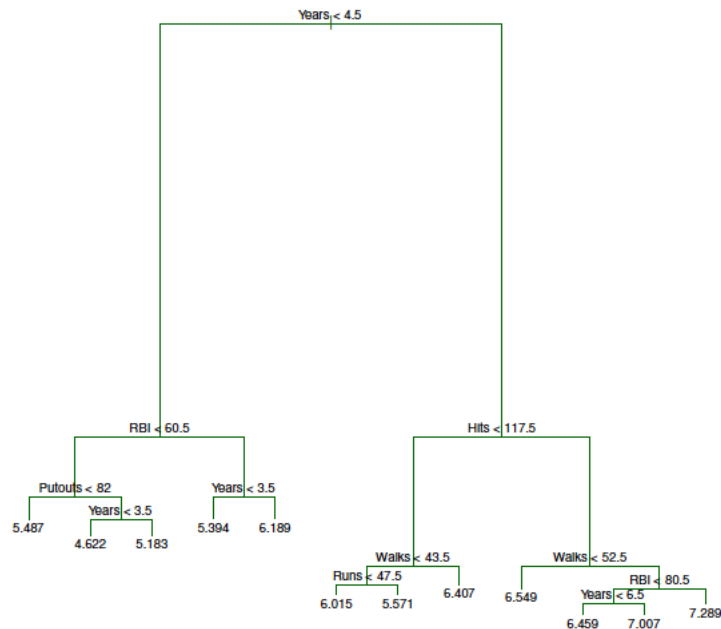
Feature importance

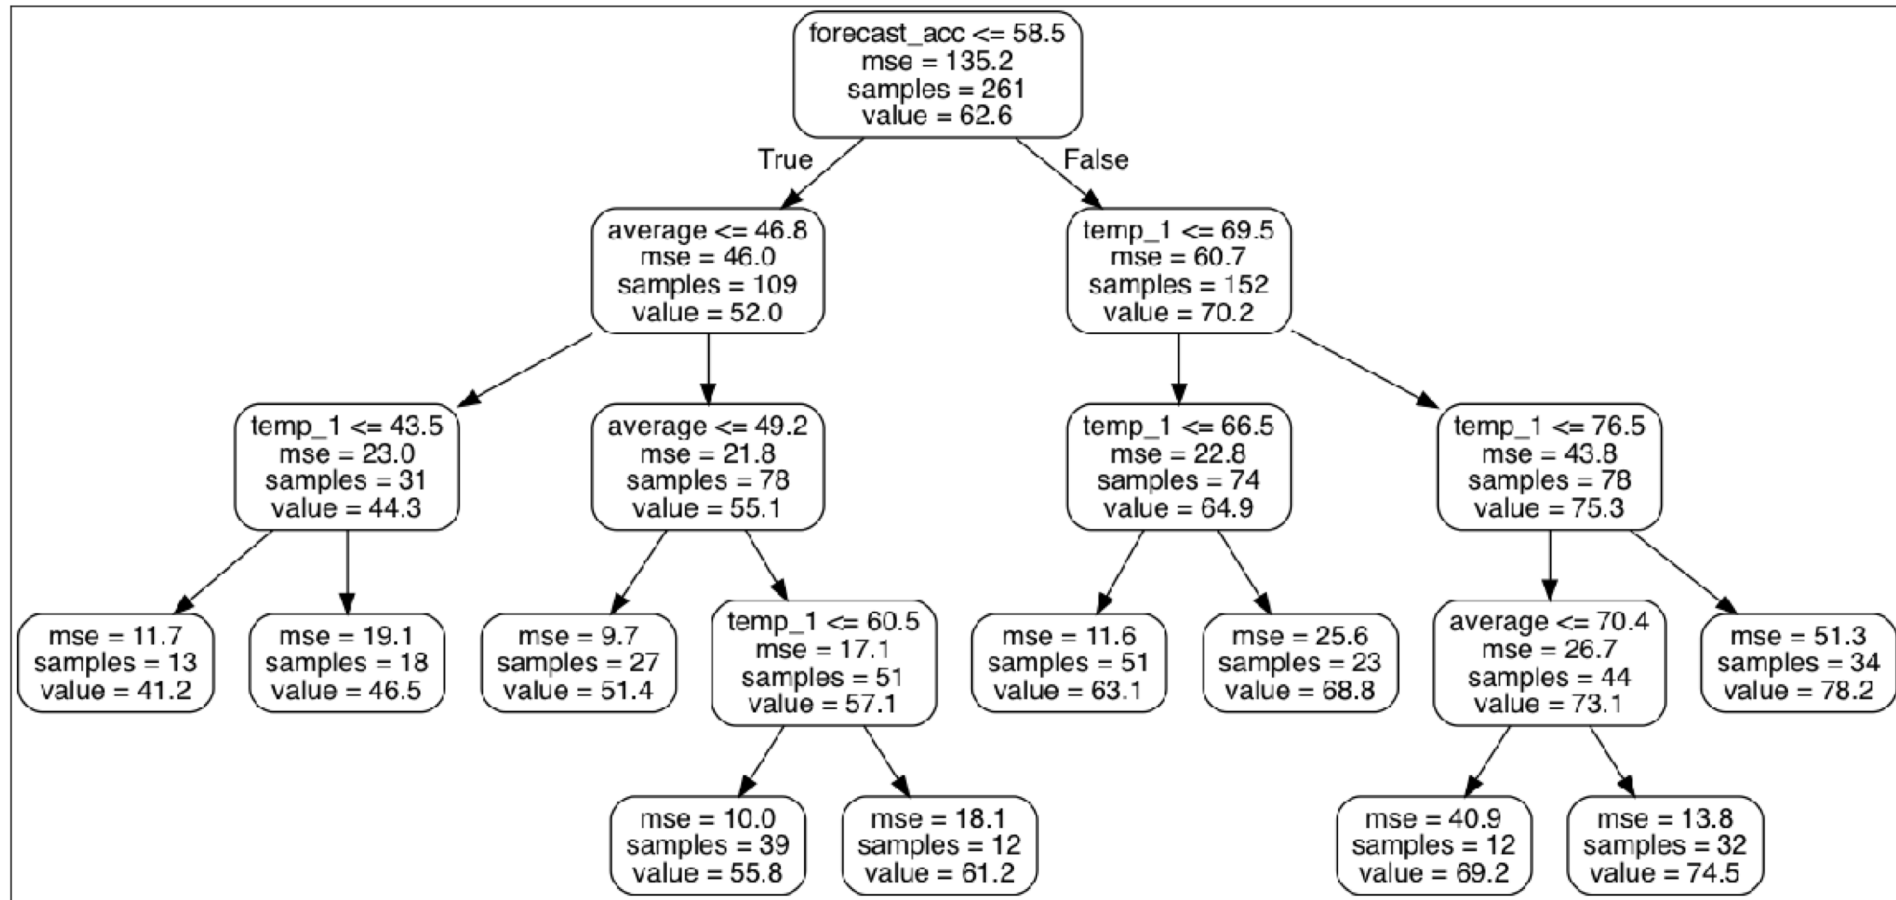❑ For each feature, find all splits where this feature was used as the split variable and add up the loss reduction at all such splits

❑ The sum reflects the importance of this feature!

Demo: weather prediction using decision tree

What about classification?

❑ The predicted class for each region = the majority class of training samples in the region

❑ How to design the tree?
- Can use the same greedy algorithm
- Split each region (by picking a feature and a threshold) to minimize a loss
- What loss functions to use?

Classification loss

❑Misclassification rate

$$L = \sum_{m=1}^{|T|} \sum_{x_n \in R_m} 1(y_n \neq \bar{y}_m) \quad , \bar{y}_m = \text{majority class of } R_m$$

❑Gini index

$$L = \sum_{m=1}^{|T|} q_m \sum_{k=1}^{K} \hat{p}_{m,k}(1 - \hat{p}_{m,k}), \ \hat{p}_{m,k} = \text{numbers of samples in class k in } R_m$$

◦ Expected error rate if we randomly pick an index, with probability $\hat{p}_{m,k}$ and error rate $1 - \hat{p}_{m,k}$

❑Cross entropy

$$L = -\sum_{m=1}^{|T|} q_m \sum_{k=1}^{K} \hat{p}_{m,k} \log \hat{p}_{m,k} \,,$$

◦ Smaller entropy means less uniform distribution (the region is more pure!)

❑Gini and cross entropy loss lead to more "Pure" regions, with 1 dominate class in each region.

Performance metric and pruning

❑ After a tree is designed, the performance is still measured by the misclassification rate or accuracy

❑ It is typical to use the Gini index or cross entropy when growing a tree, but use the misclassification rate for pruning a tree

Example: Classifying heart disease

From http://web.stanford.edu/class/stats202/content/lec19.pdf
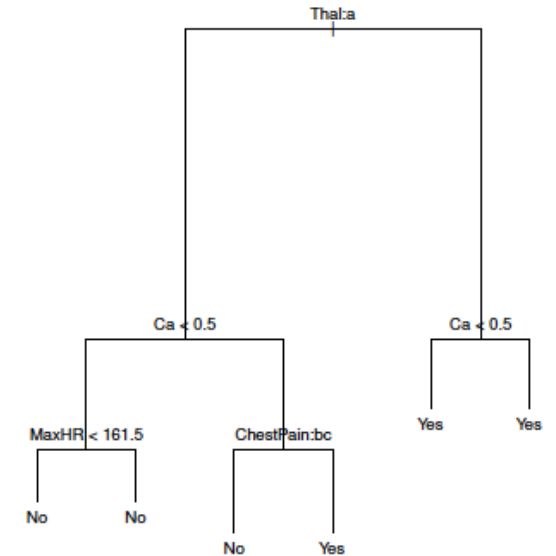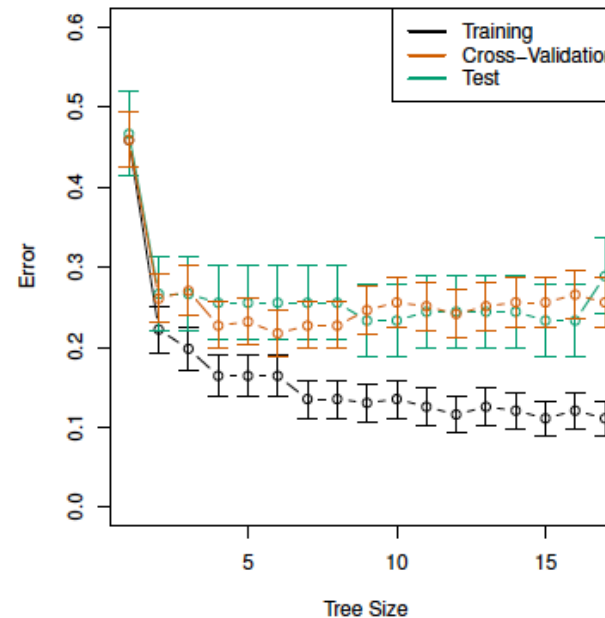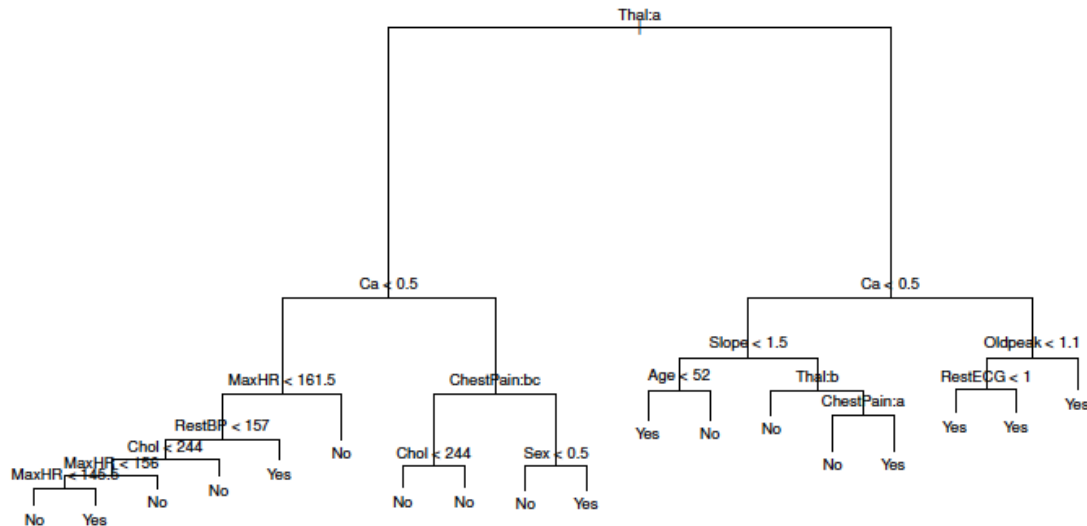
Advantage of decision tree

❑ Easy to interpret: Doctors like them

❑ Closer to human decision making

❑ Feature importance can be derived during training

❑ Can easily handle mixed type of features (numerical and categorical) and missing features in some samples
  ◦ Did not discuss here

❑ Problem:
  ◦ To reduce bias, needs to grow the tree deeper
  ◦ Deeper trees tend to overfit the training data (Large variance among different training data)
  ◦ How to overcome ?

Bagging (Bootstrap Aggregating)

❑Idea: Generate multiple trees from different training sets, and apply all models to each test sample and take average (or majority) of the results from all the trees

❑How to generate different training sets giving a dataset?

❑Cross validation: using a subset of data each time for training and the remaining for testing

❑Bootstrap sampling: Sampling by replacement, each sampling contains the same number of samples as the original dataset, but some samples are replicated, others were not included

❑Bagging: Generate B models from B bootstrap samplings
- Regression: Average the prediction results from B models
- Classification: Take the majority class index

❑Apply to other regressors/classifiers as well.

Out of bag (OOB) error

❑ Each time we draw a bootstrap sampling, we only use ~63% of the samples
- ◦ Probability that a sample is chosen among N samples in each bootstrap sampling

$$1 - \left(1 - \frac{1}{N}\right)^N \sim 1 - e^{-1} = 0.632$$

❑ We can use the remaining samples for testing

❑ OOB Error
- ◦ For each sample x_n, find the models generated by samplings which do not contain x_n. There are about 0.37B of models. Average predictions by these models for x_n.
- ◦ Compute the regression/classification error for x_n
- ◦ Average the error over all samples

❑ We can use OOB error as an estimate of the test error.

❑ Does not require design multiple models for multiple folds as in cross validation. OOB can be estimated from one pass of designing multiple trees.
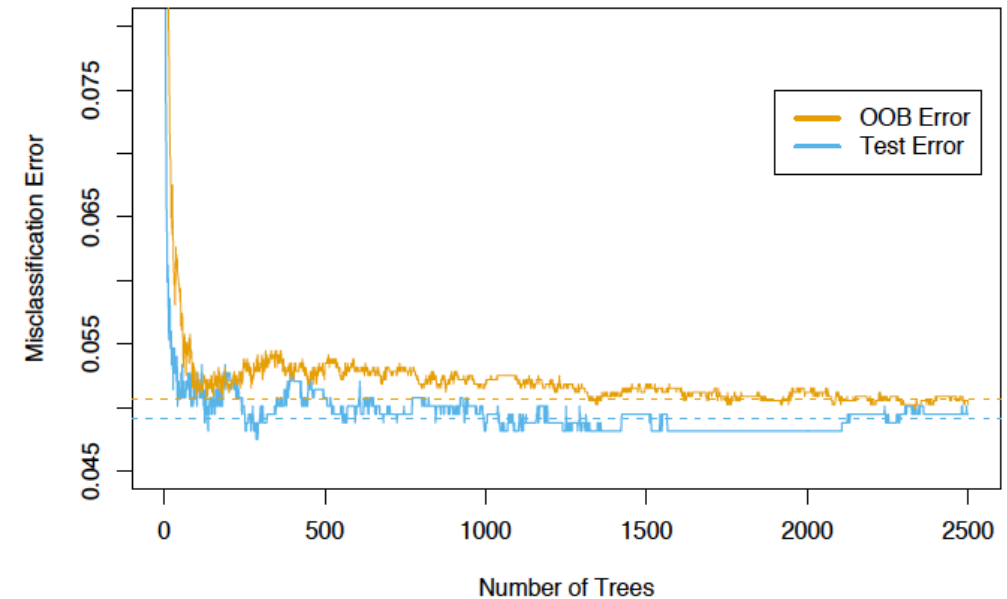
From ESL Fig. 15.4

Why bagging?

❑ When a regressor or classifier has tendency to overfit (i.e. sensitive to the training set), bagging reduces the variance of the prediction
- ◦ Reduce the test error
- ◦ Particularly useful for decision trees

❑ When the sample number N in a given dataset is large
- ◦ The empirical distribution is similar to the true distribution
- ◦ Each bootstrap sampling is similar to an independent realization of the true distribution
- ◦ Bagging amounts to averaging the fits from many identically distributed datasets

Problems with bagging?

❑ Trees generated by different samplings can be very similar

❑ Test error reduces slowly as B increases

- $f_b(x)$: prediction by tree b for test sample x
- Assume $f_b(x)$ for all b have the same mean μ and variance σ^2
- Assume these predictions have pair-wise correlation ρ
- The variance of the average prediction $f(x) = \frac{1}{B}\sum_b f_b(x)$: (Shown on board)

$$\sigma_B^2 = \rho\,\sigma^2 + \frac{1}{B}(1-\rho)\sigma^2$$

Random Forest

❑ As with Bagging: fit a different tree for each bootstrap sampling

❑ Recall that when growing a tree, at each current node (region), we split the region by choosing a particular feature and a threshold. The feature and the threshold are chosen among all P features to minimize a certain loss.

❑ With random forest, randomly choose among a subset of features (P'<P) for splitting each node

❑ The resulting trees are more different

❑ Rule of thumb: $P' = \sqrt{P}$ (but should be turned using test error or OOB error)
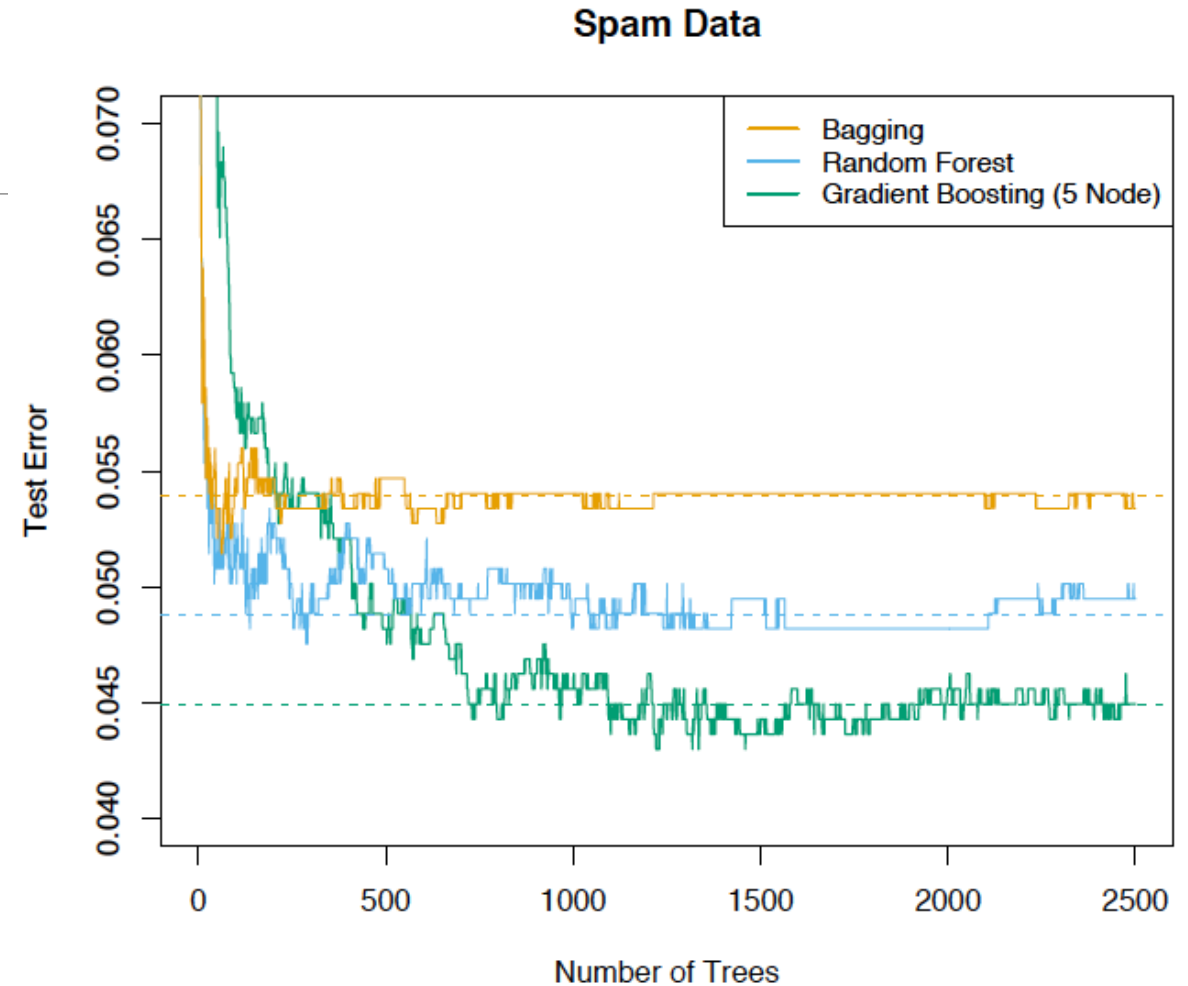
Bagging vs. RF

☐ Bagging:

$$\sigma_B^2 = \rho\,\sigma^2 + \frac{1}{B}(1-\rho)\sigma^2$$

☐ Random forest (assuming $\rho = 0$):
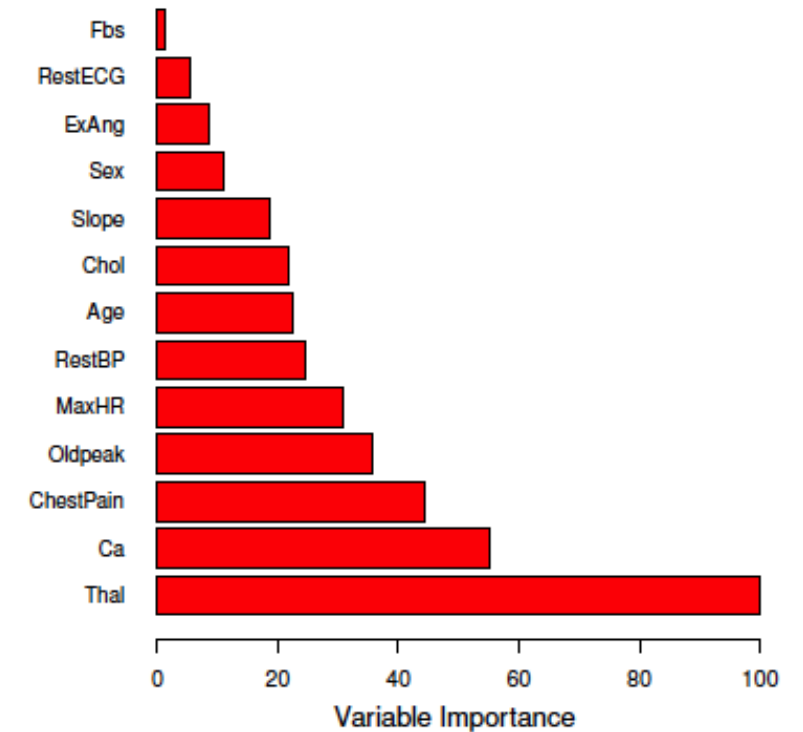
$$\sigma_B^2 = \frac{1}{B}\sigma^2$$

☐ Recall:

Test error = bias^2+ Variance +Noise Variance

Spam Data

Legend:
- Bagging
- Random Forest
- Gradient Boosting (5 Node)

(Y-axis: Test Error, X-axis: Number of Trees)

From ESL, Fig. 15.1

Feature importance

❑ For each feature, add up the loss reduction at splits where this feature was used over all trees.

Demo: Random forest

Problem with bagging and random forest

❑ Resulting model has many trees!

❑ Lose interpretability!

❑ Related methods (not covered):
  ◦ Boosting
  ◦ Gradient boosting

What you should know from this lecture

❑ How to use/interpret decision tree ?

❑ How to train a decision tree ?
  ◦ Loss function for regression
  ◦ Loss function for classification

❑ How to reduce overfitting ?

❑ What does bagging mean ?

❑ How to train and use a random forest ?

❑ How to determine feature importance?

References

❑ [ESL] Hastie, T., & Tibshirani, R. & Friedman, J.(2008). The Elements of Statistical Learning; Data Mining, Inference and Prediction. Sec. 9. (Decision tree), Sec. 15.2 (Random forest)